

Logistic Regression

Importing the data


```
% Load Data
% The first two columns contain the exam scores and the third column
% contains the label.
data = load('ex2data1.txt');
X = data(:, [1, 2]);
y = data(:, 3);
```

Visualising the data

```
plotData(X,y);

% Labels and Legend
xlabel('Exam 1 score')
ylabel('Exam 2 score')

% Specified in plot order
legend('Admitted', 'Not admitted')
```



Testing the sigmoid function

```
sigmoid(0)
```

```
ans = 0.5000
```

Cost function and gradient

```
% Setup the data matrix appropriately
[m, n] = size(X);

% Add intercept term to X
X = [ones(m, 1) X];

% Initialize the fitting parameters
initial_theta = zeros(n + 1, 1);
```

```
% Compute and display the initial cost and gradient
[cost, grad] = costFunction(initial_theta, X, y);
fprintf('Cost at initial theta (zeros): %f\n', cost);
```

Cost at initial theta (zeros): 0.693147

```
disp('Gradient at initial theta (zeros:'); disp(grad);
```

Gradient at initial theta (zeros):
-0.1000
-12.0092
-11.2628

```
% Set options for fminunc
options =
optimoptions(@fminunc, 'Algorithm', 'Quasi-Newton', 'GradObj', 'on', 'MaxIter', 400);

% Run fminunc to obtain the optimal theta
% This function will return theta and the cost
[theta, cost] = fminunc(@(t) (costFunction(t, X, y)), initial_theta, options);
```

Local minimum found.

Optimization completed because the size of the gradient is less than
the value of the optimality tolerance.

<stopping criteria details>

```
% Print theta
fprintf('Cost at theta found by fminunc: %f\n', cost);
```

Cost at theta found by fminunc: 0.203498

```
disp('theta:'); disp(theta);
```

theta:
-25.1613
0.2062
0.2015

Evaluating Logistic Regression


```
% Predict probability for a student with score 45 on exam 1
% and score 85 on exam 2
prob = sigmoid([1 45 85] * theta);
fprintf(['For a student with scores 45 and 85, we predict an ' ...
        'admission probability of %f\n\n'], prob);
```

For a student with scores 45 and 85, we predict an admission probability of 0.776291

```
% Compute accuracy on our training set
p = predict(theta, X);
fprintf('Train Accuracy: %f\n', mean(double(p == y)) * 100);
```

Train Accuracy: 89.000000

```
% Plot Boundary
plotDecisionBoundary(theta, X, y);
% Add some labels
hold on;
% Labels and Legend
xlabel('Exam 1 score')
ylabel('Exam 2 score')
% Specified in plot order
legend('Admitted', 'Not admitted')
hold off;
```




Regularized Logistic Regression

Loading the dataset

```
% The first two columns contains the X values and the third column  
% contains the label (y).  
dataset = load('ex2data2.txt');  
X = dataset(:, [1, 2]);  
y = dataset(:, 3);
```

Visualizing the data

```
plotData(X, y);  
% Put some labels  
hold on;  
% Labels and Legend  
xlabel('Microchip Test 1')  
ylabel('Microchip Test 2')  
% Specified in plot order  
legend('y = 1', 'y = 0')  
hold off;
```



Feature Mapping

```
% Add Polynomial Features  
% Note that mapFeature also adds a column of ones for us,  
% so the intercept term is handled
```

```
X = mapFeature(X(:,1), X(:,2));
```

Cost Function and gradient

```
% Initialize fitting parameters
initial_theta = zeros(size(X, 2), 1);

% Set regularization parameter lambda to 1
lambda = 1;

% Compute and display initial cost and gradient for regularized logistic regression
[cost, grad] = costFunctionReg(initial_theta, X, y, lambda);
fprintf('Cost at initial theta (zeros): %f\n', cost);
```

Cost at initial theta (zeros): 0.693147

```
% Initialize fitting parameters
initial_theta = zeros(size(X, 2), 1);

lambda = 1;
% Set Options
options =
optimoptions(@fminunc, 'Algorithm', 'Quasi-Newton', 'GradObj', 'on', 'MaxIter', 1000);

% Optimize
[theta, J, exit_flag] = fminunc(@(t)(costFunctionReg(t, X, y, lambda)), ...
    initial_theta, options);
```

Local minimum found.

Optimization completed because the size of the gradient is less than
the value of the optimality tolerance.

<stopping criteria details>

```
% Print theta
fprintf('Cost at theta found by fminunc: %f\n', cost);
```

Cost at theta found by fminunc: 0.693147

```
disp('theta:'); disp(theta);
```

```
theta:
    1.2727
    0.6252
    1.1811
   -2.0200
   -0.9174
   -1.4316
    0.1240
   -0.3655
   -0.3572
   -0.1751
   -1.4582
   -0.0510
```

```

-0.6156
-0.2747
-1.1928
-0.2422
-0.2060
-0.0447
-0.2778
-0.2954
-0.4564
-1.0432
0.0278
-0.2924
0.0156
-0.3274
-0.1439
-0.9246

```


```

% Plot Boundary
plotDecisionBoundary(theta, X, y);
hold on;
title(sprintf('lambda = %g', lambda))

% Labels and Legend
xlabel('Microchip Test 1')
ylabel('Microchip Test 2')

legend('y = 1', 'y = 0', 'Decision boundary')
hold off;

```



```
% Compute accuracy on our training set
```

```
p = predict(theta, X);
```

```
fprintf('Train Accuracy: %f\n', mean(double(p == y)) * 100);
```

```
Train Accuracy: 83.050847
```